

1 Global Constraints

We've seen at least one of these in action: all-different – now we will talk about them in a bit more.

Global constraints allow expressing a constraint between a non-fixed number of variables.

Global constraints have many advantages, including that they:

- let us express constraints that occur in our models more compactly and intuitively;
- may come with special algorithms to compute them more efficiently;
- may allow for better inference/propagation (beyond basic arc consistency).

The plan:

- look at all-different
 - increased expressiveness
 - matching-based algorithm
- look at the circuit constraint
 - sub-tour elimination via a union-find-based chain rule
- (maybe) look at other global constraints via the MiniZinc catalogue and tutorial
 - cardinality constraint

1.1 All-different

Remember our requirement from assignment 1 that all arrivals must be on different days? At least one possible solution used all-different.

n -Queens is a classic problem used to introduce all-different (and a fun problem!) so we'll discuss it.




Figure 1: Eight queens. (User:Lee Daniel Crocker – modified by Fispaul at da.wikipedia, CC BY-SA 3.0, via Wikimedia Commons.)

Instance: a natural number n ;

Question: can we place n queens on an $n \times n$ chessboard so that no two are in the same row, column, or diagonal?

Expressing this without all-different is possible but annoying. With all-different it is straightforward in MiniZinc:

(we'll look at an example board in lecture to make sure we understand encoding)

```
int: n;
array [1..n] of var 1..n: X;
include "alldifferent.mzn";
%distinct columns
constraint alldifferent(X);
% distinct diagonals upwards
constraint alldifferent([ X[i] + i - 1 | i in 1..n]);
% distinct diagonals downwards
constraint alldifferent([ X[i] - i + 1 | i in 1..n]);

solve satisfy;
```

Partial search tree from the MiniZinc tutorial:




Figure 39: Partial search trees for 9 queens: (a) the state after the addition of $q[1] = 1$, $q[2] = 4$, $q[3] = 5$ (b) the initial propagation on adding further $q[6] = 4$.

Figure 2: Partial search tree for the n -queens problem.

So hopefully I've convinced you that it's convenient for encoding. Let's talk inference (in a smaller pseudo-code example):

$$x, y, z \in \mathbb{N}, \quad D(x) = D(y) = \{1, 2\}, \quad D(z) = \{1, 2, 3\}, \quad \text{all-different}(x, y, z)$$

What can we infer about z ? Could we get this from arc consistency if we just had $x \neq y$, $x \neq z$, $y \neq z$ as our constraints?

This sort of thing might be familiar to anyone who does sudoku puzzles, and is a powerful thing

about all-different.

1.2 Matchings for all-different

We want to exclude domain values that are not part of any all-different value assignment. Luckily, graph theory can help us with **maximum matchings**.

Let's visualise our possible assignments as a bipartite graph (that is, a graph with two parts that only has edges between the parts):

- the vertex set of one part is the variables that are all-different;
- the vertex set of the other part is the possible values;
- there is an edge from variable v to value d if d is in $D(v)$.

(We will draw this out for our small example above.)

A maximum matching on a bipartite graph like this is a set of edges such that:

- no two edges share an endpoint;
- there are as many edges in this set as possible.

If our problem is satisfiable, then there must be a matching that is as large as the set of all-different variables.

Proof sketch, where V_1 is the set of variables and V_2 is the set of values:

- *Satisfiable implies matching.* Take the satisfying assignment $\Phi : V_1 \rightarrow V_2$. As it is satisfying, we know that $|\Phi| = |V_1|$ and $\Phi(v_i) \neq \Phi(v_j)$ for all distinct $v_i, v_j \in V_1$. Then take the set of edges $E_M = \{(v, x)\}$ where $v \in V_1$, $x \in V_2$, and $\Phi(v) = x$. We still need to show that:
 - $|E_M| = |V_1|$;
 - E_M are edges in our graph;
 - for every pair of edges (v, x) and (u, z) in E_M we have that $v \neq u$ and $x \neq z$.
- Let's do the second part together in lecture.

Alright, so we're only interested in edges that are in *some* max matching, but how can we tell that?

Jack Edmonds and Claude Berge can help:



Figure 3: Jack Edmonds and Claude Berge.

(Claim due to Berge, algorithm from Edmonds.)

An edge belongs to some but not all maximum matchings if and only if, for an arbitrary maximum matching, it belongs to either an even alternating path which begins at a free node, or an even alternating cycle.

OK, we've been lazy. Some definitions. Let M be a matching.

- An edge in M is a *matching edge*; every edge not in M is *free*.
- A node is *matched* if it is incident to a matching edge, and *free* otherwise.
- An *alternating path* or *cycle* is a path or cycle whose edges are alternately matching and free.

We'll draw a picture in lecture – possibly a little one and a big one.

What complexity is this? Note that you couldn't really work it out without specifying our alternating path algorithm – we'll just state it here to consider how fast it is compared with a full search tree.

If we are enforcing all-different over k variables, and the largest domain over those variables is size m , then $O(km\sqrt{k})$ to get matchings information, which dominates.

The main thing to notice is that this is much faster than a full search, so likely saves us time.

A few resources:

- <https://www.cs.upc.edu/~erodri/webpage/cps//theory/cp/global-constraints/slides.pdf>

- <https://www.sciencedirect.com/science/article/pii/S1574652606800106>
- <http://www.constraint-programming.com/people/regin/papers/globalCpaior.pdf>

1.3 Circuit constraint

We’ve just seen all-different, filtered to full arc consistency via maximum bipartite matching. Now consider a second global constraint, **circuit**, together with a filtering technique for it based on a union-find structure.

1.3.1 Definition

Fix $n \in \mathbb{N}$ and index a set of n objects (cities, tasks, locations) by $\{1, \dots, n\}$. Introduce one variable per index,

$$\text{succ}[i] \in \{1, \dots, n\}, \quad i = 1, \dots, n,$$

with the intended meaning that object i is followed by object $\text{succ}[i]$.

The constraint **circuit**(succ) holds exactly when succ, viewed as a function $\{1, \dots, n\} \rightarrow \{1, \dots, n\}$, is a single n -cycle: succ is a fixed-point-free permutation of $\{1, \dots, n\}$, and iterating succ from any index visits all n indices before returning to it.

This is the natural model for a travelling-salesperson-style problem: assign every location a successor, then (in the optimisation version) minimise total edge cost.

1.3.2 All-different is necessary but not sufficient

succ must certainly be all-different: no two indices can share a successor, and no index can be its own successor. But a fixed-point-free permutation need not be a single cycle – it can decompose into several disjoint cycles.

Example. Take $n = 4$ and $\text{succ} = [2, 1, 4, 3]$. This is a fixed-point-free permutation, so all-different is satisfied, yet

$$1 \rightarrow 2 \rightarrow 1 \quad \text{and} \quad 3 \rightarrow 4 \rightarrow 3$$

are two disjoint 2-cycles rather than one 4-cycle (Figure 4).



Figure 4: $\text{succ} = [2, 1, 4, 3]$: all-different holds, but the result is two disjoint 2-cycles, not a single circuit.

All-different filtering therefore cannot rule out sub-tours; an additional propagation rule is required.

1.3.3 Chains and a pruning rule

Suppose a subset of the `succ` variables has been fixed. The fixed edges $(i, \text{succ}[i])$ decompose into disjoint directed paths, or *chains*; an index whose variable is not yet fixed forms a chain of length one on its own.

For an index x lying on some chain, write $\text{start}(x)$ for the first index of that chain and $\text{end}(x)$ for the last, and let $\text{len}(x)$ be the number of indices on the chain. Fixing $\text{succ}[i] = j$ merges the chain ending at i with the chain starting at j ; the merged chain runs from $\text{start}(i)$ to $\text{end}(j)$, with length $\text{len}(i) + \text{len}(j)$.

Pruning rule. If x lies on a chain with $\text{start}(x) = s$, $\text{end}(x) = e$, and $\text{len}(x) < n$, remove s from $D(\text{succ}[e])$: setting $\text{succ}[e] = s$ would close a cycle of length $\text{len}(x) < n$, a sub-tour. Once $\text{len}(x) = n$, the edge $\text{succ}[e] = s$ closes the unique Hamiltonian circuit and must not be pruned.

Maintaining `start`, `end`, and `len` under repeated merges is exactly what a **union-find** (disjoint-set) structure with path compression provides, in amortised time $O(\alpha(n))$ per operation, where α is the inverse Ackermann function.

This is essentially the `nocycle` propagator of Caseau and Laburthe (1997); see the resources below.

1.3.4 Worked example

Take $n = 5$ and fix $\text{succ}[1] = 3$, then $\text{succ}[3] = 5$, then $\text{succ}[5] = 2$, in that order. Two filters act throughout: the pruning rule above, and all-different (an already-claimed successor value cannot be reused elsewhere).

Initially $D(\text{succ}[i]) = \{1, \dots, 5\} \setminus \{i\}$ for every i .

- $\text{succ}[1] = 3$: chain $1 \rightarrow 3$, length 2. Pruning rule removes 1 from $D(\text{succ}[3])$, giving $D(\text{succ}[3]) = \{2, 4, 5\}$.
- $\text{succ}[3] = 5$: chain $1 \rightarrow 3 \rightarrow 5$, length 3. Pruning rule removes 1 from $D(\text{succ}[5])$, giving $D(\text{succ}[5]) = \{2, 3, 4\}$.
- $\text{succ}[5] = 2$: chain $1 \rightarrow 3 \rightarrow 5 \rightarrow 2$, length 4. Pruning rule removes 1 from $D(\text{succ}[2])$.




Figure 5: State after fixing $\text{succ}[1] = 3$, $\text{succ}[3] = 5$, $\text{succ}[5] = 2$: a single chain $1 \rightarrow 3 \rightarrow 5 \rightarrow 2$ of length $4 < 5$, index 4 still unattached, and the closing edge $\text{succ}[2] = 1$ pruned.

At this point $D(\text{succ}[2]) = \{1, 3, 4, 5\} \setminus \{1, 3, 5\} = \{4\}$: 3 and 5 are removed by all-different (already claimed by 1 and 3 respectively), and 1 by the pruning rule. So $\text{succ}[2] = 4$ is forced.

This extends the chain to $1 \rightarrow 3 \rightarrow 5 \rightarrow 2 \rightarrow 4$, length $5 = n$. Now $D(\text{succ}[4]) = \{1, 2, 3, 5\} \setminus \{2, 3, 5\} = \{1\}$, and since $\text{len} = n$ the pruning rule permits $\text{succ}[4] = 1$.




Figure 6: The completed circuit $1 \rightarrow 3 \rightarrow 5 \rightarrow 2 \rightarrow 4 \rightarrow 1$.

Propagation alone determines the full circuit from three initial decisions, with no search.

1.3.5 Completeness

Union-find keeps this filtering cheap: $O(n\alpha(n))$ overall. It is not, however, complete. Full generalised arc consistency for **circuit** would require deciding, for every remaining candidate value of every variable, whether it extends to some Hamiltonian circuit consistent with the current domains. Taking domains directly from a graph's adjacency structure reduces this decision problem to deciding Hamiltonicity of that graph, which is NP-complete (Karp, 1972). So no polynomial-time filtering algorithm achieves full consistency for **circuit** unless $P = NP$. The pruning rule above

is sound but incomplete, and is used together with all-different's matching-based filter, which is complete for all-different on its own.

1.3.6 Circuit in MiniZinc

```
int: n;
array[1..n] of var 1..n: succ;
include "circuit.mzn";
constraint circuit(succ);

solve satisfy;
```

The all-different filtering, chain tracking, and sub-tour pruning above all take place inside this single `circuit(succ)` call.

The optimisation version (shortest tour, i.e. classic TSP) adds costs and an objective:

```
int: n;
array[1..n, 1..n] of int: cost;
array[1..n] of var 1..n: succ;
include "circuit.mzn";
constraint circuit(succ);

solve minimize sum(i in 1..n)(cost[i, succ[i]]);
```

A few resources:

- Y. Caseau and F. Laburthe, “Solving Small TSPs with Constraints,” *Proc. 14th International Conference on Logic Programming (ICLP)*, MIT Press, 1997, pp. 316–330. Available at <https://www.dcs.gla.ac.uk/~pat/cpM/papers/caseau97solving.pdf>
- R. M. Karp, “Reducibility Among Combinatorial Problems,” in *Complexity of Computer Computations*, Plenum Press, 1972, pp. 85–103.
- <https://www.minizinc.org/doc-2.3.0/en/lib-globals.html> (see `circuit`)

1.4 Global cardinality constraint

This is a generalisation of the all-different constraint:

- a set of variables X take their values from domain D , and the constraint restricts the number of times that any value can be assigned to variables, within some range.

Example: remember rota scheduling? We might have some minimum and maximum possible number of people working each night/day shift. Of course we can do this without a global constraint, but it makes it easier (and faster!) as with all-different.

Pulling a code snippet out for the guards rota example (full examples and problem description at https://github.com/magicicada/cp/tree/main/week_by_week/week_3/afternoon_examples/guard_rota):

Without global cardinality:

```
% the number of guards that must be on duty in each shift of each day
array[DAYS] of var lwbDay .. upbDay: totalOnDay;
constraint forall(d in DAYS)(
    totalOnDay[d] = sum(g in GUARDS)(roster[g,d] = DAY));
array[DAYS] of var lwbNight .. upbNight: totalOnNight;
constraint forall(d in DAYS)(
    totalOnNight[d] = sum(g in GUARDS)(roster[g,d] = NGT));
```

With:

```
% the number of guards that must be on duty in each shift of each day
array[DAYS] of var lwbDay..upbDay: totalOnDay;
array[DAYS] of var lwbNight..upbNight: totalOnNight;
constraint forall(d in DAYS)(
    global_cardinality([roster[g,d] | g in GUARDS],
        [DAY,NGT],[totalOnDay[d],totalOnNight[d]]));
```

1.5 Summary

As a general message: these special global constraints are expressive, and typically faster than rolling your own. Some, like all-different, admit filtering that is both cheap and complete; others, like circuit, only admit filtering that is cheap but incomplete (full consistency for circuit is NP-hard) – still very much worth having, just not the whole story on their own.

Try to use them when you can! See the MiniZinc global constraints library.